

A Survey of Deep Learning Methods for Document Understanding

Abstract

This paper reviews recent advances in optical character recognition and document understanding. We examine detection and recognition architectures, end-to-end training strategies, and applications to scanned academic literature. The survey covers more than fifty publications from the past decade.

1. Introduction

Scanned documents remain a major source of unstructured data. Automatic extraction of textual content from page images requires robust detection of text regions followed by accurate recognition of characters within each region. These two stages are typically trained independently.